

1. Collect Dataset of 5 leaves (Mango, Guava, Jackfruit, Neem, Tulsi) 500 each.
2. Upload in drive



3. Open co-lab and Execute step by step.

```
✓ 27s [1] from google.colab import drive
      drive.mount('/content/drive')

⇨ Mounted at /content/drive
```

4. Find dataset.

```
▶ import os

dataset_path = "/content/drive/MyDrive/LeafClassification/dataset"

# Optional: Check if the path exists
if os.path.exists(dataset_path):
    print("Dataset folder found.")
else:
    print("Dataset folder not found.")

⇨ Dataset folder found.
```

5. Split data in new folder.

```
import shutil
import os
import random
Loading...
def split_data(dataset_path, output_path, train_size=0.7, val_size=0.15, test_size=0.15):
    for folder_name in os.listdir(dataset_path):
        folder_path = os.path.join(dataset_path, folder_name)

        if os.path.isdir(folder_path):
            # List all image files in the folder
            files = os.listdir(folder_path)
            random.shuffle(files)

            # Split files into train, val, test
            num_files = len(files)
            train_split = int(train_size * num_files)
            val_split = int(val_size * num_files)

            # Create train/val/test subdirectories if they don't exist
            for split in ['train', 'val', 'test']:
                os.makedirs(os.path.join(output_path, split, folder_name), exist_ok=True)

            # Move files into respective directories
            for i, file in enumerate(files):
                src_path = os.path.join(folder_path, file)
                if i < train_split:
                    shutil.copy(src_path, os.path.join(output_path, 'train', folder_name, file))
                elif i < train_split + val_split:
                    shutil.copy(src_path, os.path.join(output_path, 'val', folder_name, file))
                else:
                    shutil.copy(src_path, os.path.join(output_path, 'test', folder_name, file))
            else:
                shutil.copy(src_path, os.path.join(output_path, 'test', folder_name, file))

# Define paths
dataset_path = "/content/drive/MyDrive/LeafClassification/dataset"
output_path = "/content/drive/MyDrive/LeafClassification/new_split"

# Split the data into train, val, and test sets
split_data(dataset_path, output_path)
```

6. View split data.

```
[ ] import os


def count_images(directory):
    for folder in os.listdir(directory):
        folder_path = os.path.join(directory, folder)
        if os.path.isdir(folder_path):
            image_count = len(os.listdir(folder_path))
            print(f"{folder}: {image_count} images")

# Define the directories
train_dir = "/content/drive/MyDrive/LeafClassification/new_split/train"
val_dir = "/content/drive/MyDrive/LeafClassification/new_split/val"
test_dir = "/content/drive/MyDrive/LeafClassification/new_split/test"

print("Training data:")
count_images(train_dir)

print("\nValidation data:")
count_images(val_dir)

print("\nTest data:")
count_images(test_dir)
```



```
Training data:
guava: 350 images
mango: 350 images
neem: 350 images
banana: 350 images
jackfruit: 350 images

Validation data:
guava: 75 images
mango: 75 images
neem: 75 images
banana: 75 images
jackfruit: 75 images

Test data:
guava: 75 images
mango: 75 images
neem: 75 images
banana: 75 images
jackfruit: 75 images
```

7. Load Data.

```
#loading the data
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Define directories for train, val, and test data
train_dir = '/content/drive/MyDrive/LeafClassification/new_split/train'
val_dir = '/content/drive/MyDrive/LeafClassification/new_split/val'
test_dir = '/content/drive/MyDrive/LeafClassification/new_split/test'

# Initialize ImageDataGenerators for data augmentation and normalization
train_datagen = ImageDataGenerator(
    rescale=1.0/255,      # Rescale images to [0, 1]
    rotation_range=20,    # Random rotations
    width_shift_range=0.2, # Random width shifts
    height_shift_range=0.2, # Random height shifts
    shear_range=0.2,      # Shear transformation
    zoom_range=0.2,       # Zoom transformation
    horizontal_flip=True,  # Horizontal flips
    fill_mode='nearest'   # Fill mode for new pixels
)

val_datagen = ImageDataGenerator(rescale=1.0/255) # Only rescale for validation

test_datagen = ImageDataGenerator(rescale=1.0/255) # Only rescale for testing

# Flow data from directories
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150), # Resize images
    batch_size=32,
    class_mode='categorical' # Multiple classes
)
```

```
    class_mode='categorical' # Multiple classes
)

val_generator = val_datagen.flow_from_directory(
    val_dir,
    target_size=(150, 150), # Resize images
    batch_size=32,
    class_mode='categorical' # Multiple classes
)

test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(150, 150), # Resize images
    batch_size=32,
    class_mode='categorical' # Multiple classes
)
```




```
Found 1749 images belonging to 5 classes.
Found 374 images belonging to 5 classes.
Found 374 images belonging to 5 classes.
```

8. Verify augmentation.

```
import matplotlib.pyplot as plt
import numpy as np

# Get a batch of augmented images from the train generator
augmented_images, _ = next(train_generator)

# Plot some augmented images
plt.figure(figsize=(10, 10))
for i in range(9): # Plot 9 images in a grid
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(augmented_images[i])
    plt.axis('off')
plt.show()
```



9. Model

```
import tensorflow as tf
from tensorflow.keras import layers, models

# Define a medium-sized CNN model
model = tf.keras.Sequential([
    # 1st Convolutional Layer
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
    layers.MaxPooling2D(2, 2),

    # 2nd Convolutional Layer
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),

    # 3rd Convolutional Layer
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),

    # 4th Convolutional Layer (Added)
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),

    # Global Average Pooling
    layers.GlobalAveragePooling2D(),

    # Fully connected layer
    layers.Dense(256, activation='relu'),

    # Output layer (5 classes: banana, guava, jackfruit, mango, neem)
    layers.Dense(5, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Display the model summary
model.summary()
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_3 (Conv2D)	(None, 15, 15, 128)	147,584
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 128)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 128)	0
dense (Dense)	(None, 256)	33,024
dense_1 (Dense)	(None, 5)	1,285

Total params: 275,141 (1.05 MB)
Trainable params: 275,141 (1.05 MB)
Non-trainable params: 0 (0.00 B)

10. Validation Data generator.

```
[ ] # Set up validation data generator (adjust this to your validation data directory)
validation_dir = '/content/drive/MyDrive/LeafClassification/new_split/val'

validation_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)

validation_generator = validation_datagen.flow_from_directory(
    validation_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical'
)
```

Found 374 images belonging to 5 classes.

11. Model fit.

```
11m history = model.fit(
    train_generator,
    epochs=10,
    validation_data=validation_generator
)
```

/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDatasetAdapter` class does not implement the `warn_if_super_not_called` method.
self.warn_if_super_not_called()

```
Epoch 1/10
55/55 ————— 139s 2s/step - accuracy: 0.4121 - loss: 1.3380 - val_accuracy: 0.6524 - val_loss: 0.7508
Epoch 2/10
55/55 ————— 123s 2s/step - accuracy: 0.7494 - loss: 0.5635 - val_accuracy: 0.7647 - val_loss: 0.4952
Epoch 3/10
55/55 ————— 120s 2s/step - accuracy: 0.7985 - loss: 0.4546 - val_accuracy: 0.7166 - val_loss: 0.6859
Epoch 4/10
55/55 ————— 142s 2s/step - accuracy: 0.8179 - loss: 0.4299 - val_accuracy: 0.7513 - val_loss: 0.5869
Epoch 5/10
55/55 ————— 121s 2s/step - accuracy: 0.8245 - loss: 0.4131 - val_accuracy: 0.8797 - val_loss: 0.2826
Epoch 6/10
55/55 ————— 150s 2s/step - accuracy: 0.8802 - loss: 0.2907 - val_accuracy: 0.8529 - val_loss: 0.3644
Epoch 7/10
55/55 ————— 118s 2s/step - accuracy: 0.8655 - loss: 0.3177 - val_accuracy: 0.9064 - val_loss: 0.2129
Epoch 8/10
55/55 ————— 118s 2s/step - accuracy: 0.8573 - loss: 0.3166 - val_accuracy: 0.8636 - val_loss: 0.3324
Epoch 9/10
55/55 ————— 119s 2s/step - accuracy: 0.8984 - loss: 0.2727 - val_accuracy: 0.9037 - val_loss: 0.1995
Epoch 10/10
55/55 ————— 119s 2s/step - accuracy: 0.9190 - loss: 0.2081 - val_accuracy: 0.9465 - val_loss: 0.1638
```

12. Save model.

```
0s model.save("/content/drive/MyDrive/LeafClassification/leaf_model_5leaf.keras")
```

13. Test.

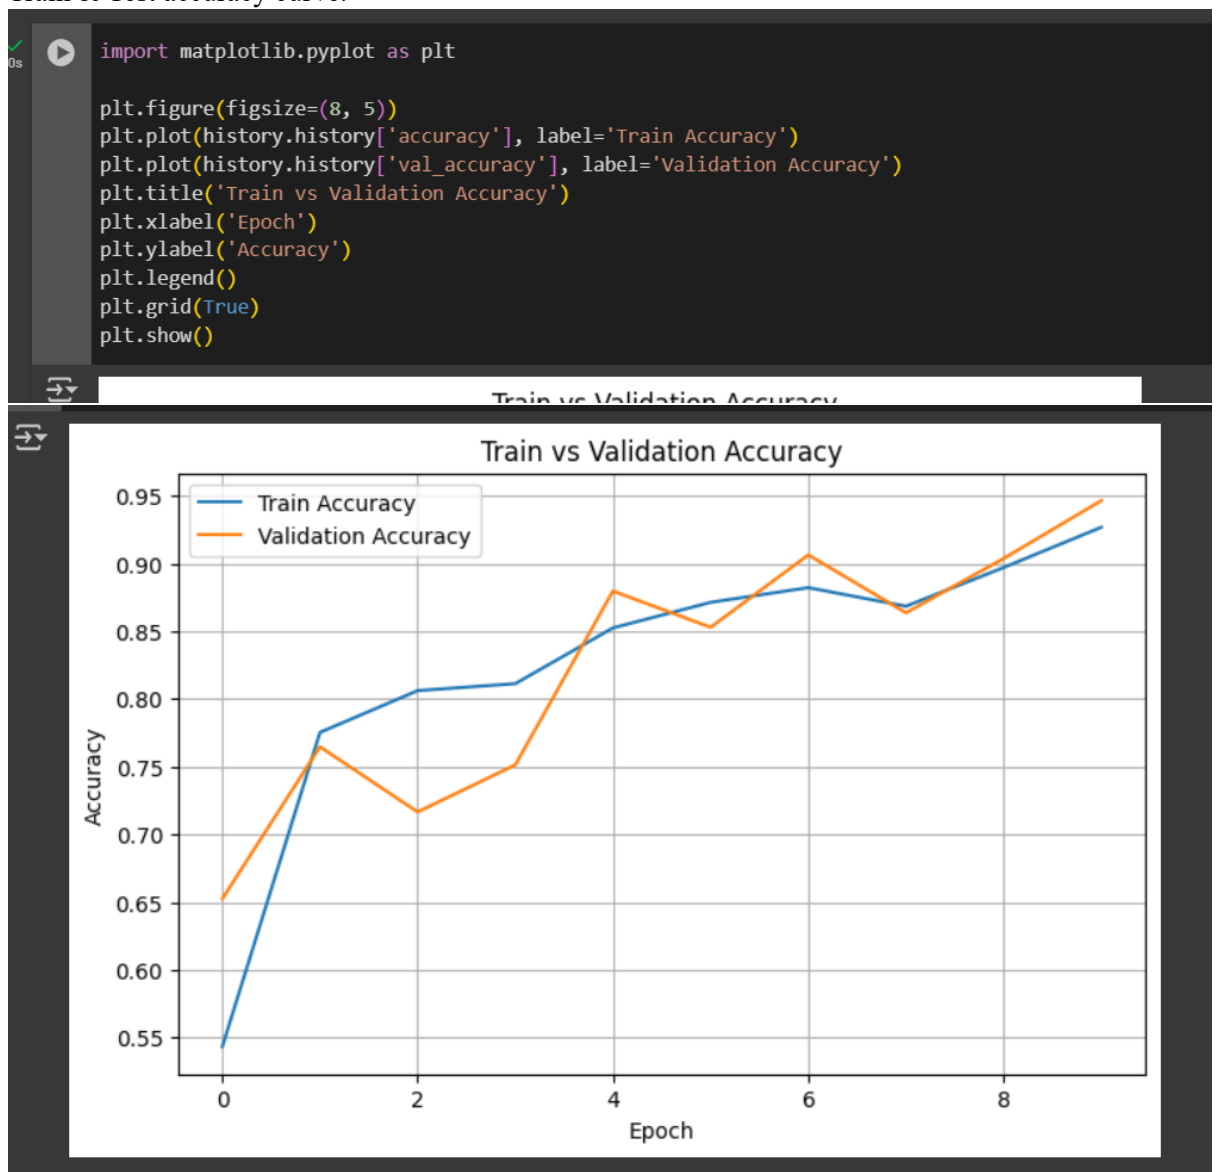
```
from tensorflow import keras

# Load the model
model = keras.models.load_model("/content/drive/MyDrive/LeafClassification/leaf_model_5leaf.keras")

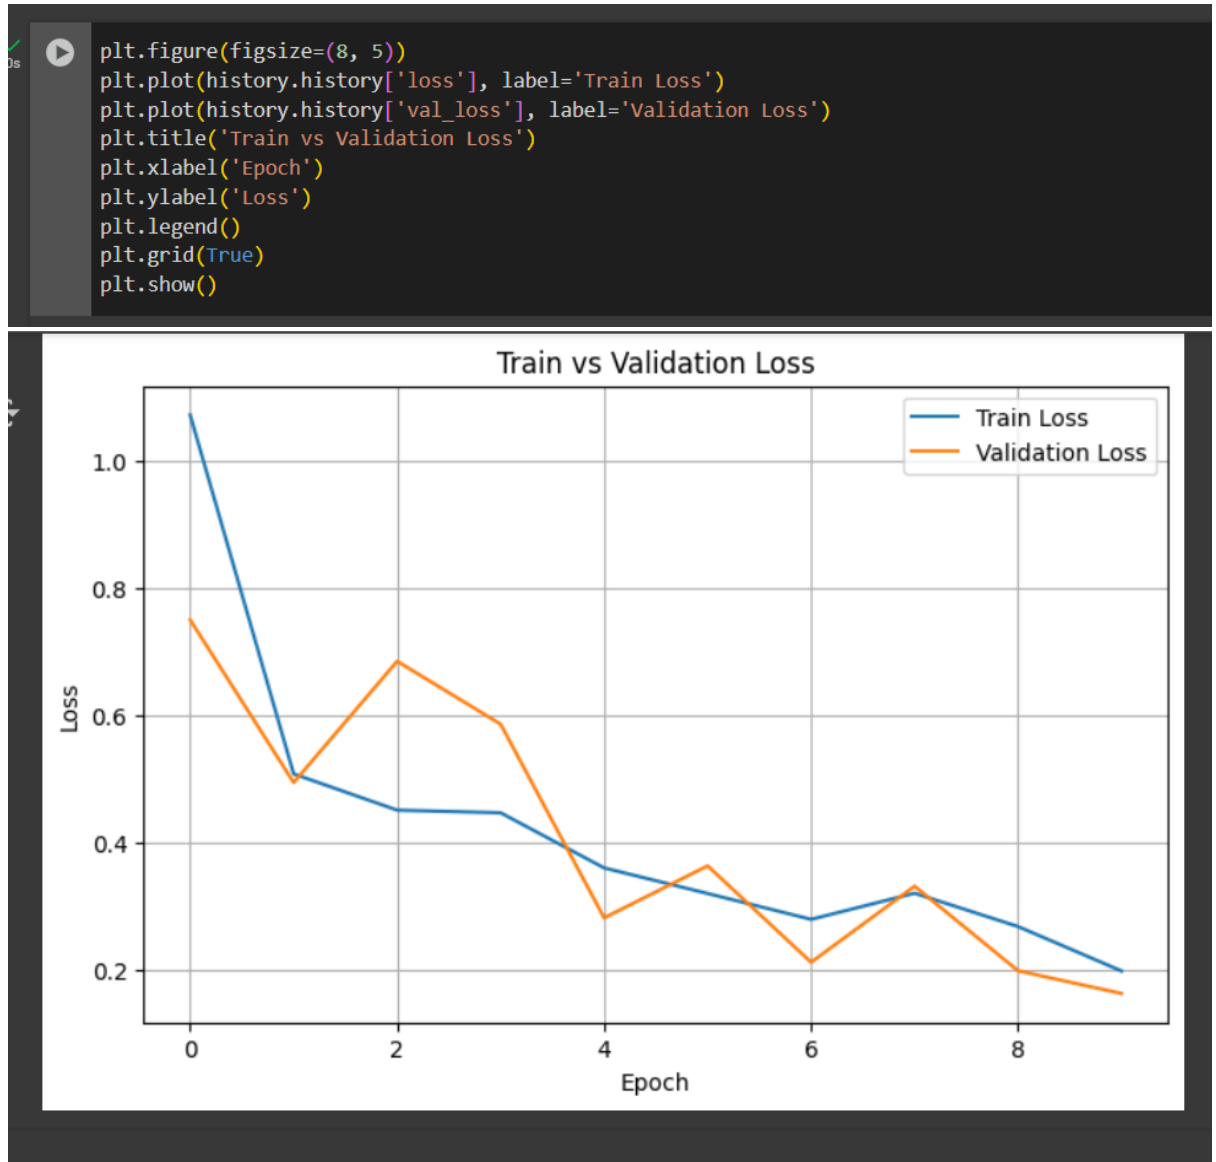
# Evaluate on test set
test_loss, test_accuracy = model.evaluate(test_generator)
print(f"Test Accuracy: {test_accuracy:.4f}")
print(f"Test Loss: {test_loss:.4f}")
```

/usr/local/lib/python3.11/dist-packages/keras/src/saving/saving_lib.py:757: UserWarning: Skipping variable saveable.load_own_variables(weights_store.get(inner_path))
12/12 ————— 9s 679ms/step - accuracy: 0.9523 - loss: 0.1481
Test Accuracy: 0.9465
Test Loss: 0.1526

14. Train & Test accuracy curve.



15. Train & Validation loss curve.



16. Test accuracy.

```
test_loss, test_accuracy = model.evaluate(test_generator)
print(f"Test Accuracy: {test_accuracy:.4f}")
```

12/12 ————— 9s 709ms/step - accuracy: 0.9343 - loss: 0.1782
Test Accuracy: 0.9465

17. Confusion Matrix.

```

import numpy as np
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

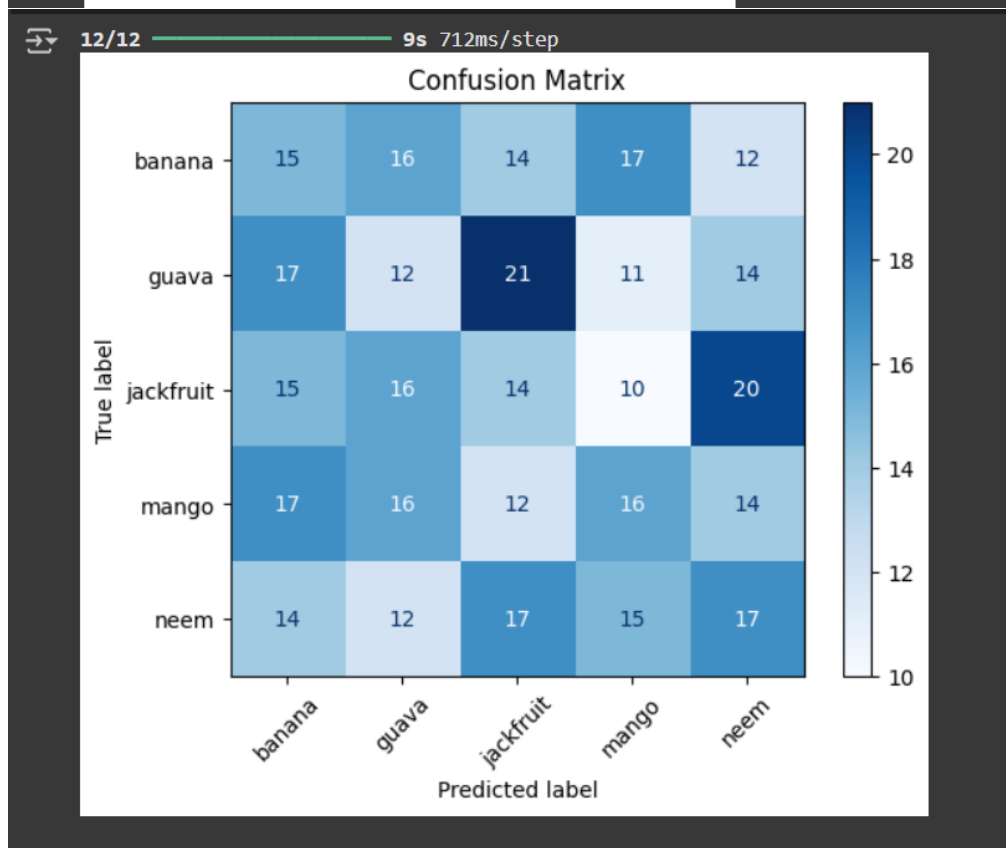
# Get ground truth labels
y_true = test_generator.classes

# Predict class probabilities
y_pred_probs = model.predict(test_generator)

# Convert probabilities to class predictions
y_pred = np.argmax(y_pred_probs, axis=1)

# Confusion matrix
cm = confusion_matrix(y_true, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=test_generator.class_indices.keys())
disp.plot(cmap='Blues', xticks_rotation=45)
plt.title('Confusion Matrix')
plt.show()

```



18. F1 Score.

```

from sklearn.metrics import f1_score

f1_macro = f1_score(y_true, y_pred, average='macro')
f1_weighted = f1_score(y_true, y_pred, average='weighted')

print(f"F1 Score (Macro): {f1_macro:.4f}")
print(f"F1 Score (Weighted): {f1_weighted:.4f}")

```

F1 Score (Macro): 0.2032
F1 Score (Weighted): 0.2031